

6.830 Project: Partitioned database with deadlock detection

Benjamin Mattocks, Wenting Zheng

May 16, 2013

Introduction

Although distributed databases have been around for a while, and many have been commercialized, there are still many interesting problems to be solved in this area. Our 6.830 project aims to implement a simple distributed database and explore some of these problems.

One of the problems we aim to tackle is data partitioning. Although some smaller distributed data may decide to replicate data across all of the machines, this will not be scalable in the long run, as the amount of data in the database increases. Therefore, *some* kind of data partitioning scheme is needed. Partitioning often brings related problems, such as hotspots (where) and . Our project aims to explore workload-based re-partitioning to increase throughput.

Another problem that is present in distributed databases is deadlock detection. Deadlocks are usually rare, but it is necessary to be able to detect them efficiently so that the database can continue to handle transactions. Distributed deadlock detection is quite hard. Since deadlock can only be detected in a global way, this makes it very easy to falsely detect deadlocks and incorrectly abort transactions.

Previous Work

TO BE DONE

Data storage

Design

Our distributed database implements a key-value storage system with transaction support. Unlike a pure key-value storage, however, it keeps the concept of tables. The tables are organized as (*primary key, record value*). The interface for accessing data contains the two operations: `write(table_name, key, value)` and `read(table_name, key)`. Thus, the data storage supports atomic groups of reads and writes.

Implementation

The implementation of the data storage is based on hash tables. Each partition contains a hash table storing the table names and a list of key-values associated with that table. Since the hashing is done on the primary keys, a partition needs to store smaller key-value stores that correspond to different tables.

Two Phase Locking and Two Phase Commit

Background

The database needs to be able to handle distributed transactions. A commit should make sure the database shows all of a transaction's changes, and an abort should make sure none of the changes are made. There can be no inconsistent state in here. Therefore, two phase locking is there to make sure that the read values are read correctly, and that the write values are written correctly and not overwritten by another transaction by the time commit happens.

On the other hand, two phase commit guarantees commits and aborts to be consistent across all machines in case of failures. The two phase commit protocol selects a coordinator for each transaction, and the coordinator ultimately decides whether that particular transaction will commit and abort after gathering information from all of its cohorts (the other servers involved in the transaction).

Implementation

Each server keeps its own lock table. The lock table provides both read locks and write locks. For this implementation, we only provide record-level locking. Transactions are only able to lock specific records. However, in the future we plan on supporting hierarchical locking.

The implementation for two-phase locking is strong strict two phase locking. Both read and write locks are held until the transaction commits or aborts. This is the implementation because the code is written in such a way that the transaction's read and write values can be requested multiple times until commit happens. Therefore, the end of phase 2 is right before commit happens, and SS2PL is the correct implementation to ensure atomic transactions.

Two phase commit requires a coordinator for each transaction. Every transaction selects a coordinator the moment it starts. Since every transaction is routed to a server to start, that server is the coordinator. The coordinator then looks into its partition table for a list of all the cohorts it needs to contact. When commit or abort happens, the coordinator contacts the necessary servers to get a list of votes, and 2PC proceeds as normal.

Our 2PC protocol is limited because the recovery process has not been written yet. Other than the logging and the recovery process, the basic 2PC algorithm is implemented. In the future we would like to look into presumed commit or presumed abort.

Partitioning

Background

Some kind of partitioning scheme is required in a distributed database. There are several ways of approaching this topic. The two most popular methods are hashing and range partitioning.

Range partitioning seems attractive when transactions tend to access records that are close to each other in a table. This often happens when transactions perform a large number of range scans. However, range partitioning requires some knowledge of the schema, and is often difficult

to implement. It is also easy to get hotspots for range partitioning.

On the other hand, hashing tends to randomize keys. This partitioning scheme is not so great for range scans since it will most likely distribute similar keys far away from each other. However, hashing does have its advantages. It is very easy to implement, which is the main reason why so many industrial implementations use hashing as its partitioning method. It is also easy to scale up as the number of keys increases.

For this project, we decided to implement a hashing-based partitioning scheme, with an adjustable number of partitions. Based on the workload and the machine limit, the user can choose to either to have fine-grained partitioning or coarse-grained partitioning.

Dynamic Re-Partitioning

Dynamic re-partitioning are often used in large distributed storage systems for a number of reasons. The main reason is to partition the data in such a way so that transactions are run with maximum efficiency. This often leads to the following:

1. Large transactions are run in parallel. For example, transactions that modify or read a large number of data sets can be executed quite fast if most of the partitions it needs to read/modify are done on separate machines in parallel.
2. Small transactions are run locally. If a transaction is very small, then it is too costly to run a distributed transaction, which would involve sending a lot of network messages.

A good and efficient partitioning scheme is very difficult design. The two previous points are somewhat contradictory with each other. Thus, in a generalized database, one would have to carefully collect transaction profiles in order to develop a good re-partitioning.

In our database we try to reduce the number of distributed transactions executed. If we have a set of small transactions that are run repeatedly, modifying and reading a small number of records, then it is better to reduce the number of distributed transactions as much as possible. The main idea of our partitioning scheme is to hash the primary keys, but finely partition the data set enough so that we can counter the randomness of hashing.

Partitions are related to each other based on how often they are run together in the same transaction. For each partition pair, we keep an *affinity factor* (or *AF*) that shows how closely related two partitions are with each other. The *AF* are increased only at the coordinator of a transaction. The reconfiguration plan is constructed based on the *AF* information

Implementation

Each server keeps track of two tables. The first table is a partition table. The partition table stores mappings of partition number to a server address. The partition table also has a single version number. The second table is an *AF* table. The *AF* table keeps a list of pairs of partitions, and a number indicating the relationship between these two partitions. Every time a successful commit happens, the coordinator of that transaction increases the *AF* number for each pair of participating partitions.

There is a server designated as the reconfiguration master. Currently, the master's identity does not change. The master server periodically sends out requests for *AF* information. Once the

master server receives all of the *AF* information, it constructs a weighted graph of all partitions in the server pool. The heavier the edge, the closer the partitions should be put next to each other.

The master then runs an algorithm to produce a reconfiguration plan. The algorithm is very simple. It simply takes the heaviest edge in the graph and brings the two partitions together on the same machine, if they are not already on the same machine. If a partition is moved from machine A to machine B, another partition must be moved from machine B to machine A. Note that the assumption here is that all partitions are almost equally involved in the workload.

Once the master decides on a plan, it sends the plan to all of the workers. The master and the workers then start the reconfiguration algorithm. The server first changes its reconfiguration state to *CHANGE*, and waits for all current transactions to commit or abort. While in this state, the server does not receive any more transactions. If a server does not have a local transaction started, or all of its transactions have been finished, the server goes to the next step, where it looks at the reconfiguration plan and sends all of the partitions that need to be sent. Because of the algorithm, most of the servers will not need to send or receive any partition. After this step is done, the server updates its partition table and increases the table's version number. It then changes the reconfiguration state to *READ*, and is ready to handle more transactions.

When there is a reconfiguration happening, transactions can come in simultaneously. It is very important to make sure that the effects of the transactions are seen by all partitions. In our database we use a very conservative approach. Transactions that involve servers that are currently reconfiguring are aborted. Transactions that involve servers with different partition tables are also aborted.

Deadlock detection

Background

A deadlock occurs when two or more process are requesting locks to items being held by other processes. More formally, it can be modeled as a Wait-For-Graph, in which a resource has a directed edge to another resource which holds a lock to an item the first resource requests. If there is a cycle in this graph, a deadlock is present.

In a distributed system, deadlock detection has been the subject of numerous studies because the distributed nature adds some complexity to the system, so certain deadlock detection algorithms might not work in a distributed environment. Distributed deadlock detection can be done with the use of centralized servers which keep track of the state of a distributed system, or they can be performed in a completely distributed manner. One problem that could occur with deadlock detection is that deadlocks might be falsely detected, which can negatively impact a system's performance. Some specific causes of this would be message delays across servers or using stale information in the WFG computed at a central server.

Implementation

We decided to implement a fully distributed algorithm, proposed by Chandy, Misra, and Haas. This is an edge-chasing algorithm, which passes probe messages between resources directly. It is based on the AND model of deadlock detection, which implements The main idea is that the origin of the first probe knows that a deadlock is present if that probe message is received by

the origin a second time. If a process is also waiting for a lock and receives a probe, it sends the probe message along to any other processes which hold locks to resources it is requesting. It retains the origin in its state so the origin will know that it sent that probe if the probe reappears there.

To implement this, we ran a specific thread for each worker on our distributed database which specifically handles these deadlock probes, independently of the main worker thread for a given transaction. This allows a worker to detect a deadlock while a specific lock is being requested. If a deadlock worker thread detects that a probe message it receives was initiated by itself, then it reports a deadlock to the worker thread and aborts the transaction. Specifically, the deadlock detection thread sets a flag in the worker thread, which then triggers an abort message to be sent by our RPC protocol to the same worker thread.

Test scenario




Figure 1: Probe message being propagated in network. The probe originating at W1 is sent to W2, and W2 sends the probe back to W1. W1 detects a deadlock upon arrival of the probe coming from W2.

We tested the simple case of two workers in separate threads wanting to obtain a lock held by the other concurrently. Suppose a worker W1 holds a lock on A and is attempting to get a lock on B. If a worker is attempting to lock a resource but cannot obtain the lock, a new probe message is sent to every other resource that holds a lock the worker requests. Suppose W2 holds the lock to B and wants the lock to A. If W1 first sends a probe message to W2, then W2 will send a probe message to W1 because W1 holds the lock to A, which W2 wants. W1 will receive this message and note that this message originated there, so it declares a deadlock and aborts itself.

Deadlock detection analysis

Our deadlock detection scheme can detect deadlocks fairly efficiently. If there are two deadlocked transactions on two servers, this can be detected by one of the transactions in a time range between 8 and 100 ms. As the number of servers increases, the detection time should scale with the number of deadlocked transactions because only deadlocked servers communicate with each other. Compared to timeout-based methods, this can be quicker and more reliable than setting an upper bound on the time spent trying to acquire a lock. In this case, it might be waiting for a transaction to complete although it's not deadlocked.

The deadlock detector runs in a separate thread, so this can cause a performance overhead. These threads are always running when a transaction is occurring, and the threads are only utilized when a possible deadlock is happening. The number of probes from an origin is no more than the number of edges in the WFG, but there may be more probes from an origin if the deadlocks are not processed in a timely manner. We did not have a specific time requirement for deadlock detection and resolution, so deadlocks might not be resolved immediately if the process obtaining a lock is still running. We set a timeout to stop the process from trying to obtain a lock and to proceed with the abort if there is in fact a detected deadlock.

Further optimizations that could occur but that we did not implement include localized deadlock detection and reduced probe frequency. That is, if the transactions in a WFG reside in the same server, the probing would not need to occur over RPC because the WFG would be local. This could save some RPC communication overhead in the real world, and because our database repartitions data according to how often values are accessed at the same time, this could present an opportunity to implement localized detection. The probe frequency could also be decreased following some optimizations provided by Chandy et al, perhaps by probing at certain time intervals.

Analysis and Evaluation

Reconfiguration Evaluation

For the reconfiguration evaluation, tests were run to compare the performance of the database with and without reconfiguration.

The tests are run with a variable number of partitions. The range of the keys are from 0 to 4000. The partitions are uniform in that the number of keys in each partition is roughly the same. The number of partitions are distributed evenly on the virtual servers. There are four threads running that make requests to the database. The four threads each represents one type of transaction. The transaction randomly reads 5 to 10 records, and randomly writes 5 to 10 records. The range of each transaction varies with each test.

Figure 2 and 3 are the data plot for the first test. The number of servers is scaled from 2 to 6. I only limited it to 6 virtual servers because we see a significant performance reduction with 8 servers due to over-using the number of cores, thus causing interference between the threads. The four transactions each access 1/4 of the database. The test is set up so that deadlock detection does not have to interfere.

As we can see, the partitioning did not cause a significant performance increase or decrease. Performance increase is not expected to be seen because the set of transactions access a random set of records from a rather large range of keys. Performance decrease is seen a little bit in figure 3, most likely due to the fact that because there are more partitions, the partitioning algorithm is moving partitions around, thus creating overhead.

The second test was run with four transactions accessing a smaller range of keys. Figure 4 and 5 show the results.

The second test also does not show a significant increase or decrease in performance. Although for 20 partitions the throughput did improve when re-partitioning was added in, 80 partitions did not seem to help as much. The reason could be that the smaller range of keys were still too big such that with 80 partitions, the re-partitioning overhead dominated the benefits.




Figure 2: Partition Test 1 – 20 partitions




Figure 3: Partition Test 1 – 80 partitions

All in all, re-partitioning does not help significantly with throughput. It seems to perform the best when the transactions' access range is not significantly greater than the partition's range. However, it can have a worse effect if the partitions are too fine. Fine partitions do not work well because of the current algorithm. Its simplicity makes it work better in the case of larger partitions. With a more sophisticated algorithm, it would make sense for fine partitions to work well. Of course, it goes without saying that transactions that access random uniformly distributed data will not see any improvement in throughput. In this case, it also did not see much of a performance degradation.




Figure 4: Partition Test 2 – 20 partitions




Figure 5: Partition Test 2 – 80 partitions

Future Work

Clearly, this project is just the beginning of a much bigger project. Many questions are still left unanswered.

Currently, two-phase commit does not have recovery process built into it. This means that the database will not endure any kind of temporary machine failure. Future work would be to add

the recovery process and logging into the database.

The database does not support complex queries such as joins and scans. This would be supported in the future. For a distributed key-value store, range scans are probably more important than actual database joins.

Re-partitioning is a very interesting problem. There are many algorithms for dynamically finding the best partitioning scheme. It is not clear whether it is important to know the global state of the relationship between partitions, or if a mere local state will be sufficient. Knowing global state would require a master to gather all of the information. It has the advantage of being very accurate, but the disadvantage of being rather slow. Local states, on the other hand, can perform ad-hoc partition transfers based on what a few servers think are the closely related partitions. The performance should be a lot faster. It is not clear whether local reconfigurations will cause the system to be very inaccurate. Even with a global set of information, the transactions' profiles may still change in such a way so that the global set is inaccurate. Future work in this area includes exploring better algorithms for global repartitioning, and find ways to repartition distributedly.

Conclusion

We implemented a distributed database that supports repartitioning and deadlock detection.