6.830 Project: Partitioned database with deadlock detection

Benjamin Mattocks, Wenting Zheng

May 16, 2013

Introduction

Our project implements a distributed database with dynamic repartioning. Repartitioning is performed to reflect the database workload such that items accessed together are placed on the same server. Our project also implements distributed deadlock detection and resolution.

Data storage

Our distributed database implements a simple key-value storage system with transaction support. The interface for accessing data contains the two operations: `write(table_name, key, value)` and `read(table_name, key)`.

Partitioning

Background

Some kind of partitioning scheme is required in a distributed database. There are several ways of approaching this topic. The two most popular methods are hashing and range partitioning.

Range partitioning seems attractive when transactions tend to access records that are close to each other in a table. This often happens when transactions perform a large number of range scans. However, range partitioning requires some knowledge of the schema, and is often difficult to implement. It is also easy to get hotspots for range partitioning.

On the other hand, hashing tends to randomize keys. This partitioning scheme is not so great for range scans since it will most likely distribute similar keys far away from each other. However, hashing does have its advantages. It is very easy to implement, which is the main reason why so many industrial implementations use hashing as its partitioning method. It is also easy to scale up as the number of keys increases.

For this project, we decided to implement a hashing-based partitioning scheme, with an adjustable number of partitions. Based on the workload and the machine limit, the user can choose to either to have fine-grained partitioning or coarse-grained partitioning.

Dynamic Re-Partitioning

Dynamic re-partitioning are often used in large distributed storage systems for a number of reasons. The main reason is to partition the data in such a way so that transactions are run with maximum efficiency. This often leads to the following:

1. Large transactions are run in parallel. For example, transactions that modify or read a large number of data sets can be executed quite fast if most of the partitions it needs to read/modify are done on separate machines in parallel.

2. Small transactions are run locally. If a transaction is very small, then it is too costly to run a distributed transaction, which would involve sending a lot of network messages.

A good and efficient partitioning scheme is very difficult design. The two previous points are somewhat contradictory with each other. Thus, in a generalized database, one would have to carefully collect transaction profiles in order to develop a good re-partitioning.

In our database we mainly tackle the second problem. If we have a set of small transactions that are run repeatedly, modifying and reading a very small number of records, then it is better to reduce the number of distributed transactions as much as possible. The main idea of our partitioning scheme is to hash the primary keys, but finely partition the data set enough so that we can counter the randomness of hashing.

Partitions are related to each other based on how often they are run together in the same transaction. For each partition pair, we keep an *affinity factor* that shows how closely related two partitions are with each other. The *affinity factors.* commits,

Deadlock detection

Background

A deadlock occurs when two or more process are requesting locks to items being held by other processes. More formally, it can be modeled as a Wait-For-Graph, in which a resource has a directed edge to another resource which holds a lock to an item the first resource requests. If there is a cycle in this graph, a deadlock is present.

In a distributed system, deadlock detection has been the subject of numerous studies because the distributed nature adds some complexity to the system, so certain deadlock detection algorithms might not work in a distributed environment. Distributed deadlock detection can be done with the use of centralized servers which keep track of the state of a distributed system, or they can be performed in a completely distributed manner. One problem that could occur with deadlock detection is that deadlocks might be falsely detected, which can negatively impact a system's performance. Some specific causes of this would be message delays across servers or using stale information in the WFG computed at a central server.

Implementation

We decided to implement a fully distributed algorithm, proposed by Chandy, Misra, and Haas. This is an edge-chasing algorithm, which passes probe messages between resources directly. It is based on the AND model of deadlock detection, which implements The main idea is that the origin of the first probe knows that a deadlock is present if that probe message is received by the origin a second time. If a process is also waiting for a lock and receives a probe, it sends the probe message along to any other processes which hold locks to resources it is requesting. It retains the origin in its state so the origin will know that it sent that probe if the probe reappears there.

To implement this, we ran a specific thread for each worker on our distributed database which specifically handles these deadlock probes, independently of the main worker thread for a given transaction. This allows a worker to detect a deadlock while a specific lock is being requested. If a deadlock worker thread detects that a probe message it receives was initiated by itself, then it reports a deadlock to the worker thread and aborts the transaction. Specifically, the deadlock detection thread sets a flag in the worker thread, which then triggers an abort message to be sent by our RPC protocol to the same worker thread.

Test scenario

Figure 1: Probe message being propagated in network. The probe originating at W1 is sent to W2, and W2 sends the probe back to W1. W1 detects a deadlock upon arrival of the probe coming from W2.

We tested the simple case of two workers in separate threads wanting to obtain a lock held by the other concurrently. Suppose a worker W1 holds a lock on A and is attempting to get a lock on B. If a worker is attempting to lock a resource but cannot obtain the lock, a new probe message is sent to every other resource that holds a lock the worker requests. Suppose W2 holds the lock to B and wants the lock to A. If W1 first sends a probe message to W2, then W2 will send a probe message to W1 because W1 holds the lock to A, which W2 wants. W1 will receive this message and note that this message originated there, so it declares a deadlock and aborts itself.

Deadlock detection analysis

Our deadlock detection scheme can detect deadlocks fairly efficiently. If there are two deadlocked transactions on two servers, this can be detected by one of the transactions in a time range between 8 and 100 ms. As the number of servers increases, the detection time should scale with the number of deadlocked transactions because only deadlocked servers communicate with each other. Compared to timeout-based methods, this can be quicker and more reliable than setting an upper bound on the time spent trying to acquire a lock. In this case, it might be waiting for a transaction to complete although it's not deadlocked.

The deadlock detector runs in a separate thread, so this can cause a performance overhead. These threads are always running when a transaction is occurring, and the threads are only utilized when a possible deadlock is happening. The number of probes from an origin is no more than the number of edges in the WFG, but there may be more probes from an origin if the deadlocks are not processed in a timely manner. We did not have a specific time requirement for deadlock detection and resolution, so deadlocks might not be resolved immediately if the process obtaining a lock is still running. We set a timeout to stop the process from trying to obtain a lock and to proceed with the abort if there is in fact a detected deadlock.

Further optimizations that could occur but that we did not implement include localized deadlock detection and reduced probe frequency. That is, if the transactions in a WFG reside in the same server, the probing would not need to occur over RPC because the WFG would be local. This could save some RPC communication overhead in the real world, and because our database repartitions data according to how often values are accessed at the same time, this could present an opportunity to implement localized detection. The probe frequency could also be decreased following some optimizations provided by Chandy et al, perhaps by probing at certain time intervals.

Analysis and Evaluation

Uniform distribution of keys:

20 partitions, 4 threads, 4000 keys; re-partitioning sleeps for 5 seconds; 1200 transactions total

1 machine: Total TPS: 3327.7499051300097

2 machines: yes par: Total TPS: 38.21019796642812 no par: Total TPS: 38.060192839969815

4 machines: yes par: Total TPS: 35.254050893555075 no par: Total TPS: 34.04016002821602

6 machines: yes par: Total TPS: 29.402536313413563 no par: Total TPS: 28.96568315101926

8 machines: yes par: Total TPS: 24.515095597585976 no par: Total TPS: 22.54833324028209

80 partitions, 4 threads, 4000 keys; re-partitioning sleeps for 5 seconds; 1200 transactions total

2 machines: yes par: Total Total TPS: 36.99719591919442 no par: Total TPS: 36.14329578690912

4 machines: yes par: Total TPS: 34.70860934996866 no par: Total TPS: 37.0490862003126

6 machines: yes par: no par:

8 machines: yes par: no par:

Future Work

Clearly, this project is just the beginning of a much bigger project. Many questions are still left unanswered.

Currently, two-phase commit does not have recovery process built into it. This means that the database will not endure any kind of

Conclusion

We implemented a distributed database that supports repartitioning and deadlock detection.